

Parallel Data Structures

Queue and Linked List

Daniel Rivera Ruiz (N10385937)

Fatima Mushtaq (N15760010)

Professor Mohamed Zahran

CSCI-GA.3033-016 Multico Processors

NYU - Courant Institute of Mathematics

Problem Definition

- Like other programming languages, the goal is to support concurrent thread-safe data-structures in **OpenMP**.
- Provide Parallel Implementation for two data structures: **Queue** and **Linked List**.

Requirements:

1. **Correctness (linearizability):** the responses received in every concurrent history are equivalent to those of a legal sequential history and their order is consistent with the real-time order.
2. **Scalability:** at least some of the operations defined for each data structure have better performance as the number of threads increases.

Experiments

For both data structure, developed two implementations: **lock-based** and **lock-free**.

- **Lock-based** implementation uses *OpenMP* lock functions: *test*, *set*, *unset*, etc.
- **Lock-free** implementation uses C++ libraries: *<atomic>*, *threads* or *OpenMP*.

Concurrent Queue Implementation supports **6 APIs**:

- *EnQueue(T), T DeQueue(), T head(), T tail(), isEmpty(), isFull()*
- *Lock free Implementation is done in a way that it uses **one Atomic Stmt only i.e.** atomic capture.*

Concurrent Linked List Implementation supports **15 APIs**, mainly:

- *add(index, T), remove(index), get(index), set(index, T), indexOf(T), contains(T)*

We performed several experiments for each implementation, varying the **number of threads** and the **type and number of operations** performed. The main metric we used to evaluate performance is execution time as measured with **omp_get_wtime()**.

Results

Concurrent Queue comparison: Enqueue vs Dequeue

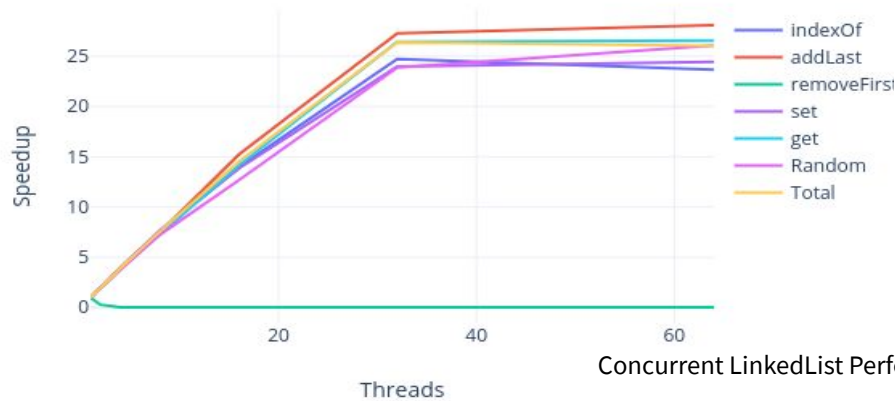


Concurrent Queue Performance

Scalability Analysis Concurrent Queue OpenMP Locks vs Atomic/LockFree

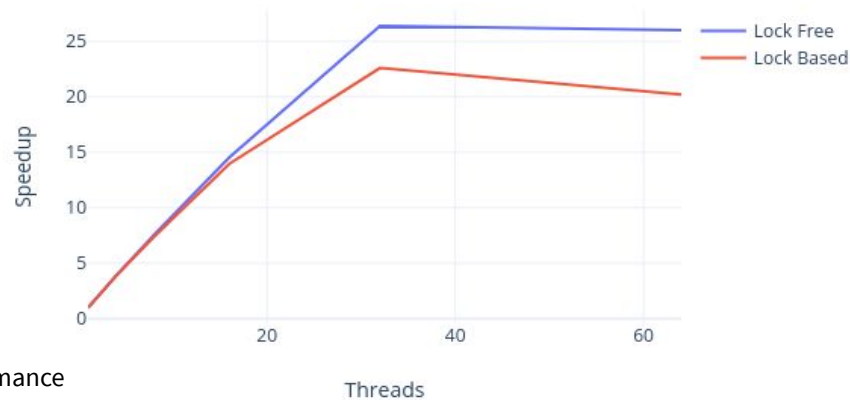


Speedup for different operations over 1,000 repetitions



Concurrent LinkedList Performance

Overall speedup for the two implementations



Conclusions

Queue:

- Pthreads/Mutex has shown to be more efficient than OpenMP Implementation.
- Lock-Free approach is complex. It showed less performance gain in OpenMP.
- Concurrent queue implementation in OpenMP scales with problem size and number of threads.

Linked List:

- Operations that try to modify the list at the same location every time (e.g. addFirst(), deleteFirst(), etc.) do not benefit from parallelism.
- The most challenging operation in the parallel implementation of a Linked List is deletion, since it requires two atomic operations to ensure correctness.